



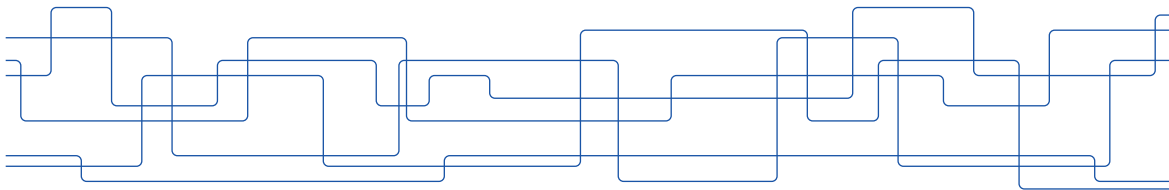
DD2434 Machine Learning, Advanced Course

Module 8 : randomization techniques in machine learning

Aristides Gionis

argioni@kth.se

KTH Royal Institute of Technology



in previous lectures

- ▶ discussed common dimensionality-reduction methods, i.e., PCA, MDS, and Isomap
- ▶ methods rely on minimizing the reconstruction error
- ▶ a typical statement : find a mapping $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$, with $k \ll d$ to minimize

$$E_f = \mathbb{E}[\|\mathbf{x} - \mathbf{y}\| - \|f(\mathbf{x}) - f(\mathbf{y})\|]$$

- ▶ so, reconstruction error guarantees distance preservation on expectation
- ▶ for a given pair of points (outliers) the error can be very large
- ▶ question : can we devise a mapping to preserve distances in the worst case?
- ▶ yes! the Johnson-Lindenstrauss lemma

overview of module 8

- ▶ essential probability tools for large-scale data analysis
- ▶ the Johnson-Lindenstrauss lemma
- ▶ data streams and computation of frequency moments

reading material

- ▶ your favorite book on probability, computing, and randomized algorithms
 - e.g., Motwani and Raghavan, Randomized algorithms (chapters 3 and 4)
- ▶ Dasgupta and Gupta. An Elementary Proof of a Theorem of Johnson and Lindenstrauss. 2002
- ▶ Achlioptas. Database-friendly Random Projections. 2003
- ▶ Alon, Matias, and Szegedy. The space complexity of approximating the frequency moments. 1999

essential probability tools

- ▶ union bound
- ▶ linearity of expectation
- ▶ concentration inequalities
 - Markov inequality, Chebyshev inequality, Chernoff bound

the union bound

- ▶ by the probability axioms we know that for any finite (or countably infinite) sequence of pairwise mutually disjoint events E_1, E_2, \dots it is

$$\Pr \left[\bigcup_{i \geq 1} E_i \right] = \sum_{i \geq 1} \Pr[E_i]$$

- ▶ moreover, for any events E_1, E_2, \dots, E_n

$$\Pr \left[\bigcup_{i=1}^n E_i \right] \leq \sum_{i=1}^n \Pr[E_i]$$

how to apply the union bound

- ▶ consider a random process for which we can identify the possible “bad” events
- ▶ assume that “bad” event i happens with probability p_i
- ▶ union bound says that probability that any “bad” event happens is at most $\sum_i p_i$
- ▶ if we can show that $\sum_i p_i$ is (significantly) less than 1
- ▶ then, probability of success (no “bad” event) is at least $1 - \sum_i p_i$

random variable

- ▶ a random variable X on a sample space Ω is a function $X : \Omega \rightarrow \mathbb{R}$
- ▶ a **discrete** random variable takes only a finite (or countably infinite) number of values

expectation and variance of a random variable

- ▶ the expectation of a discrete random variable X , denoted by $\mathbb{E}[X]$, is given by

$$\mathbb{E}[X] = \sum_x x \mathbb{P}[X = x] = \mu_X$$

where the summation is over all values in the range of X

- ▶ variance

$$\text{Var}[X] = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[(X - \mu_X)^2] = \sigma_X^2$$

linearity of expectation

- ▶ for **any** two random variables X and Y

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$$

- ▶ for a constant c and a random variable X

$$\mathbb{E}[cX] = c \mathbb{E}[X]$$

linearity of expectations — application to coupon collector

- ▶ consider a collector of coupons (e.g., stamps, coins, football cards, etc.)
- ▶ assume n different coupon types
- ▶ in each trial, the collector picks a coupon type at random
- ▶ how many trials are needed, in expectation, until the collector gets all the coupon types?

analysis (1/2)

- ▶ let c_1, c_2, \dots, c_X the sequence of coupon types picked, $c_i \in \{1, \dots, n\}$
- ▶ so, the random variable X denotes the total number of trials until all types are collected
- ▶ call c_i success if a new coupon type is picked
 - c_1 and c_X are always successes
- ▶ divide the sequence in epochs:
 - the i -th epoch starts after the i -th success and ends with the $(i+1)$ -th success
 - thus, i ranges from 0 to $n-1$
- ▶ define X_i to be the length of the i -th epoch
- ▶ clearly $X = \sum_{i=0}^{n-1} X_i$

analysis (2/2)

- ▶ probability of success in the i -th epoch

$$p_i = \frac{n-i}{n}$$

- ▶ X_i is geometrically distributed with parameter p_i , and its expectation is

$$E[X_i] = \frac{1}{p_i} = \frac{n}{n-i}$$


- ▶ from **linearity of expectation**

$$E[X] = E\left[\sum_{i=0}^{n-1} X_i\right] = \sum_{i=0}^{n-1} E[X_i] = \sum_{i=0}^{n-1} \frac{n}{n-i} = n \sum_{i=1}^n \frac{1}{i} = nH_n$$

where H_n is the harmonic number, asymptotically equal to $\ln n$

concentration inequalities

- ▶ also known as tail inequalities, or concentration bounds
- ▶ we want to bound the probability that a random variable deviates from its expectation
- ▶ many different bounds, depending on the kind of distribution we are interested in



Markov inequality

Theorem

let X a random variable taking non-negative values

for all $t > 0$

$$\Pr[X \geq t] \leq \frac{E[X]}{t}$$

or equivalently

$$\Pr[X \geq k E[X]] \leq \frac{1}{k}$$

- ▶ one of the simplest bounds, it makes no assumption other than non-negativity
- ▶ however, because of its generality, it is also a weak bound

Proof.

- ▶ by the definition of expectation $\mathbb{E}[f(X)] = \sum_x f(x) \mathbb{P}[X = x]$
- ▶ define $f(x) = 1$ if $x \geq t$ and 0 otherwise
- ▶ applying the first equation with this particular function f gives $\mathbb{E}[f(X)] = \mathbb{P}[X \geq t]$
- ▶ notice also that $f(x) \leq x/t$, which implies

$$\mathbb{E}[f(X)] \leq \mathbb{E}\left[\frac{X}{t}\right]$$

- ▶ putting everything together

$$\mathbb{P}[X \geq t] = \mathbb{E}[f(X)] \leq \mathbb{E}\left[\frac{X}{t}\right] = \frac{\mathbb{E}[X]}{t}$$

the last step is linearity of expectation



Chebyshev inequality

Theorem

let X a random variable with expectation μ_X and standard deviation σ_X

then for all $t > 0$

$$\Pr[|X - \mu_X| \geq t\sigma_X] \leq \frac{1}{t^2}$$

Proof.

- ▶ notice that

$$\mathbb{P}[|X - \mu_X| \geq t\sigma_X] = \mathbb{P}[(X - \mu_X)^2 \geq t^2\sigma_X^2]$$

- ▶ the random variable $Y = (X - \mu_X)^2$ has expectation σ_X^2
- ▶ apply the Markov inequality on Y



Chernoff bounds

Theorem

let X_1, \dots, X_n independent Poisson trials, i.e., $\mathbb{P}[X_i = 1] = p_i$ and $\mathbb{P}[X_i = 0] = 1 - p_i$

define $X = \sum_i X_i$, so $\mu = \mathbb{E}[X] = \sum_i \mathbb{E}[X_i] = \sum_i p_i$

for any $\delta > 0$

$$\mathbb{P}[X > (1 + \delta)\mu] \leq e^{-\frac{\delta^2 \mu}{3}}$$

and

$$\mathbb{P}[X < (1 - \delta)\mu] \leq e^{-\frac{\delta^2 \mu}{2}}$$

a simple example of using the Chernoff bound

- ▶ consider n coin flips; define $X_i = 1$ if i -th coin flip is H and 0 if T
- ▶ the total number of H's is $X = \sum_i X_i$, and expectation is $\mu = E[X] = n/2$
- ▶ pick $\delta = \frac{2c\sqrt{n}}{n}$, so $(1 - \delta)\mu = (1 - \frac{2c\sqrt{n}}{n})\frac{n}{2} = \frac{n}{2} - c\sqrt{n}$
- ▶ for that δ the r.h.s. of the Chernoff bound will be $e^{-\frac{\delta^2\mu}{2}} = e^{-\frac{4c^2 \cdot n \cdot n}{n^2 \cdot 2 \cdot 2}} = e^{-c^2}$ which drops very fast with c
- ▶ the Chernoff bound gives

$$\mathbb{P}\left[X < \frac{n}{2} - c\sqrt{n}\right] = \mathbb{P}[X < (1 - \delta)\mu] \leq e^{-\frac{\delta^2\mu}{2}} = e^{-c^2}$$

and similarly $\mathbb{P}\left[X > \frac{n}{2} + c\sqrt{n}\right] \leq e^{-\frac{\delta^2\mu}{3}} = e^{-2c^2/3}$

- ▶ so, probability that number of H's is outside the range $\left[\frac{n}{2} - c\sqrt{n}, \frac{n}{2} + c\sqrt{n}\right]$ is very small

- ▶ there is a extensive amount of work dedicated to concentration inequalities
 - distribution of the random variable
 - additive vs. multiplicative deviations
 - independent vs. dependent settings
- ▶ useful in randomized algorithms, large-scale data analysis, machine learning theory

what we want to achieve in a nutshell

- ▶ given a set of n points $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ in \mathbb{R}^d we want to construct a mapping $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$, with $k \ll d$, so that

$$\|\mathbf{x}_i - \mathbf{x}_j\| \approx \|f(\mathbf{x}_i) - f(\mathbf{x}_j)\|, \quad \text{for all } \mathbf{x}_i, \mathbf{x}_j \in \mathcal{X}$$

- ▶ the emphasis here is on the “for all”
 - notice that methods like PCA preserve distances “on expectation”
- ▶ it is also important that we aim for $k \ll d$
 - the result implies that we can embed a set of points in a much lower dimensional space without losing much information
- ▶ the mapping f will be linear, and will be constructed using **random projections**
- ▶ this result has many applications in machine learning and theoretical computer science

the Johnson-Lindenstrauss lemma, precise formulation

Theorem (Johnson-Lindenstrauss, 1984)

for any $0 < \epsilon < 1$ and any integer n , let

$$k \geq \frac{4 \ln(n)}{\epsilon^2/2 - \epsilon^3/3}$$

then for any set of points $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ in \mathbb{R}^d there exists a mapping $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$, such that for all $\mathbf{x}_i, \mathbf{x}_j \in \mathcal{X}$:

$$(1 - \epsilon) \|\mathbf{x}_i - \mathbf{x}_j\|^2 \leq \|f(\mathbf{x}_i) - f(\mathbf{x}_j)\|^2 \leq (1 + \epsilon) \|\mathbf{x}_i - \mathbf{x}_j\|^2$$

some observations about the Johnson-Lindenstrauss lemma

- ▶ result is about existence of a mapping
 - however, such a mapping can be constructed very easily
- ▶ ϵ is a parameter that controls the quality of distance preservation
- ▶ dimensionality k of lower-dimensional space grows with $\mathcal{O}(\epsilon^{-2})$
 - this is intuitive, as smaller distortion requires larger dimension of the host space
- ▶ dimensionality k of lower-dimensional space grows with $\ln(n)$ (number of points)
- ▶ d can be as large as n
 - (but not larger, as n points lie always on a $(n - 1)$ -dimensional hyperplane)
- ▶ thus, the reduction in dimension can be exponentially large (from n to $\mathcal{O}(\ln(n))$)
 - this can be very useful for algorithms that are exponential in d
 - after a random projection such algorithms become polynomial in d
- ▶ on the other hand, the lemma is useful only if the dimension d is large enough
 - in particular, it should be $d = \omega(\ln(n))$

random projections

1-dimensional space

- ▶ let \mathbf{z} be a random vector drawn from the uniform distribution on the unit sphere in \mathbb{R}^d
- ▶ the function $\pi_{\mathbf{z}} : \mathbb{R}^d \rightarrow \mathbb{R}$ with $\pi_{\mathbf{z}}(\mathbf{x}) = \langle \mathbf{z}, \mathbf{x} \rangle = \mathbf{z}^T \mathbf{x}$, for $\mathbf{x} \in \mathbb{R}^d$ is a random projection of \mathbf{x} on the 1-dimensional space

k -dimensional space

- ▶ let $\mathbf{z}_1, \dots, \mathbf{z}_k$ be k random vectors drawn from the uniform distribution on the unit sphere in \mathbb{R}^d
- ▶ let \mathcal{Z} be the subspace spanned by $\{\mathbf{z}_1, \dots, \mathbf{z}_k\}$, and let \mathbf{Z} be a $k \times d$ matrix whose rows are the vectors $\mathbf{z}_1, \dots, \mathbf{z}_k$
- ▶ the function $\pi_{\mathcal{Z}} : \mathbb{R}^d \rightarrow \mathbb{R}^k$ with $\pi_{\mathcal{Z}}(\mathbf{x}) = (\langle \mathbf{z}_1, \mathbf{x} \rangle, \dots, \langle \mathbf{z}_k, \mathbf{x} \rangle) = \mathbf{Z} \mathbf{x}$, for $\mathbf{x} \in \mathbb{R}^d$ is a random projection of \mathbf{x} on the k -dimensional space

expected length of random projections

Proposition (1-dimensional random projection)

let \mathbf{x} be a fixed vector in \mathbb{R}^d with $\|\mathbf{x}\| = 1$; let $\pi_{\mathbf{z}}(\mathbf{x})$ be its projection on a random vector \mathbf{z} sampled from the uniform distribution on the unit sphere in \mathbb{R}^d ; then

$$\mathbb{E}_{\mathbf{z}} [|\pi_{\mathbf{z}}(\mathbf{x})|^2] = 1/d$$

- (random projection shrinks vector lengths by $\sqrt{1/d}$)

Proof.

► by a symmetry argument we can assume that $\mathbf{x} = \mathbf{e}_1 = (1, 0, \dots, 0)$

► then

$$\mathbb{E}_{\mathbf{z}} [|\pi_{\mathbf{z}}(\mathbf{x})|^2] = \mathbb{E}_{\mathbf{z}} [|\pi_{\mathbf{z}}(\mathbf{e}_1)|^2] = \mathbb{E}_{\mathbf{z}} [|\langle \mathbf{z}, \mathbf{e}_1 \rangle|^2] = \mathbb{E}_{\mathbf{z}} [|z_1|^2]$$

► the random vector \mathbf{z} has unit length

$$1 = \|\mathbf{z}\|^2 = \mathbb{E}_{\mathbf{z}} [\|\mathbf{z}\|^2] = \mathbb{E}_{\mathbf{z}} \left[\sum_i |z_i|^2 \right] = \sum_i \mathbb{E}_{\mathbf{z}} [|z_i|^2]$$

► by symmetry again

$$\mathbb{E}_{\mathbf{z}} [|z_i|^2] = 1/d \quad \text{for all } i = 1, \dots, d$$

► we can conclude

$$\mathbb{E}_{\mathbf{z}} [|\pi_{\mathbf{z}}(\mathbf{x})|^2] = \mathbb{E}_{\mathbf{z}} [|z_1|^2] = 1/d$$

□

expected length of random projections

Proposition (k -dimensional random projection)

let \mathbf{x} be a fixed vector in \mathbb{R}^d with $\|\mathbf{x}\| = 1$; let $\pi_{\mathcal{Z}}(\mathbf{x})$ be its projection on a k -dimensional random space \mathcal{Z} ; then

$$\mathbb{E}_{\mathcal{Z}}[\|\pi_{\mathcal{Z}}(\mathbf{x})\|^2] = k/d$$

- (random projection on k -dimensional space shrinks vector lengths by $\sqrt{k/d}$)

Proof.

- ▶ consider the rotation \mathbf{R} that maps \mathcal{Z} to the k -dimensional space \mathcal{E}_k spanned by the k basis vectors $\{\mathbf{e}_1, \dots, \mathbf{e}_k\}$, that is, $\mathcal{E}_k = \mathbf{R}\mathcal{Z}$; also define $\mathbf{y} = \mathbf{R}\mathbf{x}$
- ▶ since a rotation does not change the vector norms we have

$$\mathbb{E}_{\mathcal{Z}} [\|\pi_{\mathcal{Z}}(\mathbf{x})\|^2] = \mathbb{E}_{\mathbf{R}\mathcal{Z}} [\|\pi_{\mathbf{R}\mathcal{Z}}(\mathbf{R}\mathbf{x})\|^2] = \mathbb{E}_{\mathbf{R}\mathcal{Z}} [\|\pi_{\mathcal{E}_k}(\mathbf{y})\|^2]$$

- ▶ length of a **fixed** unit vector projected on a **random** subspace =
= length of a **random** unit vector projected on a **fixed** subspace
- ▶ therefore, we have

$$\mathbb{E}_{\mathbf{R}\mathcal{Z}} [\|\pi_{\mathcal{E}_k}(\mathbf{y})\|^2] = \mathbb{E}_{\mathbf{y}} [\|\pi_{\mathcal{E}_k}(\mathbf{y})\|^2] = \mathbb{E}_{\mathbf{y}} \left[\sum_{i=1}^k |\langle \mathbf{e}_i, \mathbf{y} \rangle|^2 \right] = \sum_{i=1}^k \mathbb{E}_{\mathbf{y}} [|y_i|^2] = k/d$$

□

so far

- ▶ we have shown that the expected length of a unit vector on a random k -dimensional subspace is $\sqrt{k/d}$
- ▶ we also want to show that the distribution is sharply concentrated around its mean

concentration properties of random projections

Proposition

let $L = \|\pi_{\mathcal{Z}}(\mathbf{x})\|^2$, be the squared length of a random projection of unit vector \mathbf{x}

then $\mathbb{E}[L] = k/d$, and

for $\beta > 1$

$$\mathbb{P}\left[L \geq \beta \frac{k}{d}\right] \leq \exp\left(\frac{k}{2}(1 - \beta + \ln \beta)\right)$$

for $\beta < 1$

$$\mathbb{P}\left[L \leq \beta \frac{k}{d}\right] \leq \exp\left(\frac{k}{2}(1 - \beta + \ln \beta)\right)$$

- (the probability that $\|\pi_{\mathcal{Z}}(\mathbf{x})\|^2$ deviates by more than a factor β from its expectation is exponentially small)

implication of concentration inequalities

Proposition

let $0 < \epsilon < 1$ and $k \geq 4(\epsilon^2/2 - \epsilon^3/3)^{-1} \ln(n)$

let \mathbf{x} be any vector, and $L = \|\pi_{\mathcal{Z}}(\mathbf{x})\|^2$

define $\mu = k/d\|\mathbf{x}\|^2$, so $\mathbb{E}[L] = \mu$

then

$$\mathbb{P}[L \geq (1 + \epsilon)\mu] \leq \frac{1}{n^2} \quad \text{and} \quad \mathbb{P}[L \leq (1 - \epsilon)\mu] \leq \frac{1}{n^2}$$

Proof.

- by the concentration inequality claimed before, with $\beta = 1 + \epsilon > 1$

$$\begin{aligned}\mathbb{P}[L \geq (1 + \epsilon)\mu] &\leq \exp\left(\frac{k}{2}(1 - (1 + \epsilon) + \ln(1 + \epsilon))\right) \\ &\leq \exp\left(\frac{k}{2}(-\epsilon + (\epsilon - \epsilon^2/2 + \epsilon^3/3))\right) = \exp\left(-\frac{k(\epsilon^2/2 - \epsilon^3/3)}{2}\right) \\ &\leq \exp(-2 \ln n) = \frac{1}{n^2}\end{aligned}$$

where the second inequality follows by $\ln(1 + x) \leq x - x^2/2 + x^3/3$,

which holds for all $x \geq 0$

and the last inequality holds by our choice of $k \geq 4(\epsilon^2/2 - \epsilon^3/3)^{-1} \ln(n)$

- the case of $\mathbb{P}[L \leq (1 - \epsilon)\mu] \leq \frac{1}{n^2}$ can be shown in a similar manner

□

- ▶ consider mapping $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$ so that $f(\mathbf{x}) = \sqrt{d/k} \pi_{\mathcal{Z}}(\mathbf{x})$
- ▶ fix two vectors \mathbf{x}_i and \mathbf{x}_j in the set $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
- ▶ consider the difference of \mathbf{x}_i and \mathbf{x}_j as a new vector $\mathbf{y} = \mathbf{x}_i - \mathbf{x}_j$
- ▶ note that $\|f(\mathbf{x}_i) - f(\mathbf{x}_j)\| = \|f(\mathbf{x}_i - \mathbf{x}_j)\| = \|f(\mathbf{y})\|$
- ▶ the expected length of $f(\mathbf{y})$ is simply $\mathbb{E}[\|f(\mathbf{y})\|^2] = \|\mathbf{y}\|^2$ and by the previous proposition

$$\mathbb{P}[\|f(\mathbf{y})\|^2 \leq (1 - \epsilon)\|\mathbf{y}\|^2] \leq \frac{1}{n^2} \quad \text{and} \quad \mathbb{P}[\|f(\mathbf{y})\|^2 \geq (1 + \epsilon)\|\mathbf{y}\|^2] \leq \frac{1}{n^2}$$

so

$$(1 - \epsilon)\|\mathbf{x}_i - \mathbf{x}_j\|^2 \leq \|f(\mathbf{x}_i) - f(\mathbf{x}_j)\|^2 \leq (1 + \epsilon)\|\mathbf{x}_i - \mathbf{x}_j\|^2 \quad (\text{JL})$$

does not hold with probability at most $2/n^2$ (union bound)

- ▶ consider now all $\binom{n}{2}$ pairs of points in \mathcal{X}
- ▶ (JL) does not hold for at least one pair with probability at most (union bound, again)

$$\frac{(n-1)n}{2} \cdot \frac{2}{n^2} = 1 - \frac{1}{n}$$

- ▶ it follows that (JL) holds for all pairs with probability at least $\frac{1}{n}$
- ▶ probability of existence of a mapping f that satisfies (JL) for all pairs is at least $\frac{1}{n}$
- ▶ non zero probability implies that such a mapping exists
 - (an instance of the probabilistic method)



how can we find such a random projection?

- ▶ we know that a “good” random projection exists with probability at least $1/n$
- ▶ we can easily sample a random projection and check if it satisfies the desired property
 - if not, we discard it and repeat
- ▶ we will need $\mathcal{O}(n)$ trials to find one with the desired property
- ▶ checking one trial requires time $\mathcal{O}(n^2 d)$
- ▶ so, we have a randomized algorithm
- ▶ total running time $\mathcal{O}(n^3 d)$, on expectation
 - quite expensive, but at least polynomial

how to construct a random projection?

- ▶ we need to sample random vectors from the uniform distribution on the unit sphere in \mathbb{R}^d
- ▶ how to sample one such a vector?

answer:

1. sample each coordinate **independently** from the **normal distribution** $\mathcal{N}(0, 1)$
2. normalize the vector to unit norm

how to construct a random projection?

- ▶ an alternative elegant construction was given by Achlioptas (2003) :
- ▶ create a $k \times d$ matrix \mathbf{Z} , where each entry z_{ij} is sampled independently as follows

$$z_{ij} = \begin{cases} 1 & \text{with probability } 1/6 \\ 0 & \text{with probability } 4/6 \\ -1 & \text{with probability } 1/6 \end{cases}$$

- ▶ the matrix \mathbf{Z} is used to define the random projection $f(\mathbf{x}) = \sqrt{3/k} \mathbf{Z} \mathbf{x}$

simplified version of Johnson-Lindenstrauss lemma

Theorem (Achlioptas, 2003)

for $\epsilon, \beta > 0$, and integer n , take

$$k \geq \frac{4 + 2\beta}{\epsilon^2/2 - \epsilon^3/3} \ln(n)$$

for any set of points $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ in \mathbb{R}^d construct the mapping f as shown in the previous slide;

then

$$(1 - \epsilon) \|\mathbf{x}_i - \mathbf{x}_j\|^2 \leq \|f(\mathbf{x}_i) - f(\mathbf{x}_j)\|^2 \leq (1 + \epsilon) \|\mathbf{x}_i - \mathbf{x}_j\|^2$$

holds for all pairs $\mathbf{x}_i, \mathbf{x}_j \in \mathcal{X}$, with probability at least $1 - \frac{1}{n^\beta}$

simplified version of Johnson-Lindenstrauss lemma

- ▶ notes on the JL version by Achlioptas
 - spherical symmetry is not essential; concentration is more crucial
 - construction succeeds with high probability; one sampled matrix is sufficient (whp)
 - matrix \mathbf{Z} is sparse, matrix-vector multiplication very easy; “database friendly”

embeddings in algorithmic design and machine learning

- ▶ the Johnson-Lindenstrauss lemma can be used to speed up many algorithms, simply by reducing the dimensionality of the data
- ▶ the topic of embedding is very broad and has many different “flavors”
 - high to low dimensional spaces
 - general metrics to vector spaces
 - graphs to trees
 - graphs to vector spaces

embeddings in algorithmic design and machine learning

- ▶ k -means clustering of d -dimensional points :
 - there exist a random projection of the data to dimension $\mathcal{O}(k/\epsilon^2)$ that preserves the k -means clustering solution to factor $2 + \epsilon$
- ▶ column subset selection problem
 - select k columns of a matrix that give the best rank- k approximation of the matrix
 - random projections can be used to give a good approximation to this problem
- ▶ streaming computation
 - compute a “sketch” over a data stream to estimate useful properties of the stream
 - many streaming algorithms rely on random projections

data streams

- ▶ a data stream is a massive sequence of data
 - too large to store (on disk, memory, cache, etc.)
- ▶ examples:
 - social media (e.g., twitter feed, foursquare checkins)
 - sensor networks (weather, radars, cameras, etc.)
 - network traffic (trajectories, source/destination pairs)
 - satellite data feed
- ▶ how to deal with such data? what are the issues?

issues when working with data streams

► space

- data size is very large
- often not possible to store the whole dataset
- inspect each data item, make some computations, not possible to store it, never get to inspect it again
- some times data is stored, but making one single pass takes a lot of time, especially when the data is stored on disk
- other times, we can afford a small number of passes over the data

► time


- data “flies by” at a high speed
- computation time per data item needs to be small

data streams

- ▶ data items can be of **complex types**
 - documents, images, geo-located time-series . . .
- ▶ to study basic algorithmic ideas we abstract away application-specific details
- ▶ consider the data stream as a **sequence of numbers**

data-stream model

for simplicity assume that the input is a stream of numbers



data-stream model

- ▶ **stream**: m elements from universe of size n , e.g.,

$$\langle x_1, x_2, \dots, x_m \rangle = 6, 1, 7, 4, 9, 1, 5, 1, 5, \dots$$

- ▶ **goal**: compute a function over the elements of the stream, e.g., median, number of distinct elements, quantiles, etc.
- ▶ **constraints**:
 1. limited working memory, sublinear in n and m , e.g., $\mathcal{O}(\log n + \log m)$,
 2. access data sequentially
 4. limited number of passes, in some cases only one pass
 4. process each element quickly, e.g., $\mathcal{O}(1)$, $\mathcal{O}(\log n)$, etc.

warm up: computing some simple functions

- ▶ assume that a number can be stored in $\mathcal{O}(\log n)$ space
- ▶ min and max can be computed with $\mathcal{O}(\log n)$ space
- ▶ sum and mean need $\mathcal{O}(\log n + \log m)$ space

$$\mu_X = \mathbb{E}[X] = \mathbb{E}[x_1, \dots, x_m] = \frac{1}{m} \sum_{i=1}^m x_i$$

- ▶ variance

$$\text{Var}[X] = \text{Var}[x_1, \dots, x_m] = \mathbb{E}[(X - \mathbb{E}[X])^2] = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_X)^2$$

- ▶ can be computed in a single pass?

how to tackle massive data streams?

- ▶ a general and powerful technique: **sampling**
- ▶ idea:
 1. keep a random sample of the data stream
 2. perform the computation on the sample
 3. extrapolate
- ▶ example: compute the median of a data stream
 - the median of the sample is a good estimate of the median on the whole stream

how to tackle massive data streams?

- ▶ a general and powerful technique: **sampling**
- ▶ idea:
 1. keep a random sample of the data stream
 2. perform the computation on the sample
 3. extrapolate
- ▶ example: compute the median of a data stream
 - the median of the sample is a good estimate of the median on the whole stream
- ▶ but ... how to keep a random sample of a data stream?

reservoir sampling

- ▶ **problem** : take a uniform sample s from a stream of unknown length while keeping in memory a single number

reservoir sampling

- ▶ **problem** : take a uniform sample s from a stream of unknown length while keeping in memory a single number
- ▶ **algorithm** :
 1. initially $s \leftarrow x_1$
 2. on seeing the t -th element, $s \leftarrow x_t$ with probability $1/t$

reservoir sampling

- ▶ **problem** : take a uniform sample s from a stream of unknown length while keeping in memory a single number
- ▶ **algorithm** :
 1. initially $s \leftarrow x_1$
 2. on seeing the t -th element, $s \leftarrow x_t$ with probability $1/t$
- ▶ **analysis** :
 - what is the probability that $s = x_i$ at some time $t \geq i$?

$$\begin{aligned}\mathbb{P}[s = x_i] &= \frac{1}{i} \left(1 - \frac{1}{i+1}\right) \cdots \left(1 - \frac{1}{t-1}\right) \left(1 - \frac{1}{t}\right) \\ &= \frac{1}{i} \frac{i}{i+1} \cdots \frac{t-2}{t-1} \frac{t-1}{t} = \frac{1}{t}\end{aligned}$$

- ▶ reservoir sampling algorithm uses $\mathcal{O}(\log n)$ bits of memory
- ▶ can easily be extended to taking k samples with $\mathcal{O}(k \log n)$ bits of memory

how to tackle massive data streams?

- ▶ another powerful technique: sketching
- ▶ idea:
 - apply a random projection that maps high-dimensional data to a small “sketch”
 - post-process sketch to estimate quantities of interest

computing frequency moments on data streams

- ▶ $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
each element x_j has a “key” in the set $N = \{1, \dots, n\}$
 $m_i = |\{j : x_j = i\}|$ the number of occurrences of key i

computing frequency moments on data streams

- ▶ $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
each element x_j has a “key” in the set $N = \{1, \dots, n\}$
 $m_i = |\{j : x_j = i\}|$ the number of occurrences of key i
- ▶ define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

computing frequency moments on data streams

- ▶ $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
each element x_j has a “key” in the set $N = \{1, \dots, n\}$
 $m_i = |\{j : x_j = i\}|$ the number of occurrences of key i
- ▶ define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

- F_0 is the number of distinct elements
- F_1 is the length of the sequence
- F_2 is the second moment: index of homogeneity, size of self-join, other applications
- F_∞^* frequency of most frequent element

computing frequency moments on data streams

- ▶ we can compute all frequency moments using $O(n \log m)$ memory bits in a straightforward manner
- ▶ can be done more efficiently?
- ▶ problem studied by Alon, Matias, and Szegedy in their seminal paper
- ▶ the idea is to create a sketch that requires small space and provides an estimate of F_k
- ▶ sketch can be considered a **random projection** on the streaming setting

estimating F_2

- recall our problem setting :

$X = (x_1, x_2, \dots, x_m)$ a sequence of elements

each element x_j has a “key” in the set $N = \{1, \dots, n\}$

$m_i = |\{j : x_j = i\}|$ the number of occurrences of key i

$$F_k = \sum_{i=1}^n m_i^k$$

estimating F_2

- recall our problem setting :

$X = (x_1, x_2, \dots, x_m)$ a sequence of elements

each element x_j has a “key” in the set $N = \{1, \dots, n\}$

$m_i = |\{j : x_j = i\}|$ the number of occurrences of key i

$$F_k = \sum_{i=1}^n m_i^k$$

- the sketching algorithm :

- hash each key $i \in \{1, \dots, n\}$ to a random $\epsilon_i \in \{-1, +1\}$
- maintain sketch $Z = \sum_i \epsilon_i m_i$, need only $\mathcal{O}(\log n + \log m)$ space
- take $X = Z^2$
- return $Y = \frac{1}{k} \sum_{j=1}^k X_j$, i.e., the average of k such estimates X_1, \dots, X_k ,
where $k = \frac{16}{\lambda^2}$, and where λ controls the accuracy of the estimate

expectation of the estimate is correct

$$\begin{aligned}\mathbb{E}[X] &= \mathbb{E}[Z^2] \\ &= \mathbb{E}\left[\left(\sum_{i=1}^n \epsilon_i m_i\right)^2\right] \\ &= \sum_{i=1}^n m_i^2 \mathbb{E}[\epsilon_i^2] + 2 \sum_{i < j} m_i m_j \mathbb{E}[\epsilon_i] \mathbb{E}[\epsilon_j] \\ &= \sum_{i=1}^n m_i^2 = F_2\end{aligned}$$

accuracy of the estimate

easy to show

$$\mathbb{E}[X^2] = \sum_{i=1}^n m_i^4 + 6 \sum_{i < j} m_i^2 m_j^2$$

accuracy of the estimate

easy to show

$$\mathbb{E}[X^2] = \sum_{i=1}^n m_i^4 + 6 \sum_{i < j} m_i^2 m_j^2$$

which gives

$$\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 = 4 \sum_{i < j} m_i^2 m_j^2 \leq 2F_2^2$$

accuracy of the estimate

easy to show

$$\mathbb{E}[X^2] = \sum_{i=1}^n m_i^4 + 6 \sum_{i < j} m_i^2 m_j^2$$

which gives

$$\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 = 4 \sum_{i < j} m_i^2 m_j^2 \leq 2F_2^2$$

and by Chebyshev's inequality

$$\Pr[|Y - F_2| \geq \lambda F_2] \leq \frac{\text{Var}[Y]}{\lambda^2 F_2^2} = \frac{\text{Var}[X]/k}{\lambda^2 F_2^2} \leq \frac{2F_2^2/k}{\lambda^2 F_2^2} = \frac{2}{k\lambda^2} = \frac{1}{8}$$

estimation of F_2

Theorem (Alon, Matias, Szegedy, 1999)

let X_1, \dots, X_k be AMS sketches, with $k = \frac{16}{\lambda^2}$, and Y be their average $Y = \frac{1}{k} \sum_{j=1}^k X_j$

then, Y is an unbiased estimator of F_2 , and the quality of the approximation is given by

$$\Pr[|Y - F_2| \geq \lambda F_2] \leq \frac{1}{8}$$

summary and discussion

- ▶ random projections have many applications, also in streaming computation
- ▶ streaming is a widely-researched topic; it has been studied in the context of
 - estimating number of distinct items in data streams
 - estimating frequencies and finding most frequent items
 - matrix approximations and PCA
 - graph mining