

# DD2418 Language Engineering: 3a: Statistical language models


Johan Boye, KTH

# Statistical properties of language


We will view words as *random events* generated by some *random process*.

Some high-level observations:

**Zipf's law:** The number of occurrences of a word is inversely proportional to the word's rank in a frequency table.



**Heap's law:** The number of unique words in a text is proportional to the square root of the number of tokens.



# Statistical language model

A (statistical) **language model** predicts the probability of the next word, based on the preceding words.

- *He wrote a \_\_\_\_\_.*
- *He wrote a letter \_\_\_\_\_.*
- *He wrote a letter to \_\_\_\_\_.*
- *He wrote a letter to his \_\_\_\_\_.*

# Applications of language models

- Word prediction

- *I'm going \_\_\_\_\_*

↙

to
for
on

- Spelling correction

- *Flights form Boston.*

- Speech recognition

- $P(\text{"recognize speech"}) > P(\text{"wreck a nice beach"})$

- Translation

- $P(\text{"tall building"}) > P(\text{"high building"})$

- ... and many more

# Some notation

$\mathbf{P}(\mathbf{w})$  = the probability of the word  $w$ .

- Picking a random word from a text, what is the probability that that word will be  $w$ ?

$\mathbf{P}(\mathbf{w}_1 \mathbf{w}_2 \dots \mathbf{w}_n)$  = the probability of the sequence  $w_1 w_2 \dots w_n$ .

- Picking a random sequence of  $i$  words from a text, what is the probability that that sequence will be  $w_1 w_2 \dots w_n$ ?

# Chain rule for probabilities

The *chain rule* rewrites a joint probability into a product of conditional probabilities.

$$P(A, B) = P(B|A)P(A)$$

In general:

$$P(A_1, \dots, A_n) = \prod_{i=1}^n P(A_i | A_1, \dots, A_{i-1})$$

For example:

$$P(\text{I really like ants}) = P(\text{ants} | \text{I really like}) P(\text{like} | \text{I really}) P(\text{really} | \text{I}) P(\text{I})$$

We take  $P(\text{ants} | \text{I really like})$  to mean

$$P(w_i = \text{"ants"} | w_{i-1} = \text{"like"}, w_{i-2} = \text{"really"}, w_{i-3} = \text{"I"})$$

# Markov assumption

Assume that a word only depends on the previous couple of words.

Unigram model: Don't consider the context.

$$P(\text{I have a unicorn}) = P(\text{I})P(\text{have})P(\text{a})P(\text{unicorn})$$

Bigram model: Look at the previous word.

$$\begin{aligned} P(\text{I have a unicorn}) = \\ P(\text{I})P(\text{have}|\text{I})P(\text{a}|\text{have})P(\text{unicorn}|\text{a}) \end{aligned}$$

In general: An *n-gram model* takes the  $n - 1$  preceding words into account. Typically such models are estimated from a large corpus.

$$P(w_i | w_1, \dots, w_{i-1}) = P(w_i | w_{i-n+1}, \dots, w_{i-1})$$

# DD2418 Language Engineering: 3b: Estimating n-gram models

Johan Boye, KTH



# Markov assumption

Assume that a word only depends on the previous couple of words.

Unigram model: Don't consider the context.

$$P(\text{I have a unicorn}) = P(\text{I})P(\text{have})P(\text{a})P(\text{unicorn})$$

Bigram model: Look at the previous word.

$$\begin{aligned} P(\text{I have a unicorn}) = \\ P(\text{I})P(\text{have}|\text{I})P(\text{a}|\text{have})P(\text{unicorn}|\text{a}) \end{aligned}$$

In general: An *n-gram model* takes the  $n - 1$  preceding words into account. Typically such models are estimated from a large corpus.

$$P(w_i | w_1, \dots, w_{i-1}) = P(w_i | w_{i-n+1}, \dots, w_{i-1})$$

# Maximum likelihood estimation (MLE)

How do we estimate n-gram probabilities  $P(w_i | w_{i-n+1}, \dots, w_{i-1})$ ?

Count word strings in a text corpus, compute fraction:

$$P(w_i | w_{i-n+1}, \dots, w_{i-1}) = \frac{c(w_{i-n+1}, \dots, w_{i-1}, w_i)}{c(w_{i-n+1}, \dots, w_{i-1})}$$

In particular (bigram probabilities):

$$P(\text{to} | \text{like}) = \frac{c(\text{like to})}{c(\text{like})}$$

Unigram probabilities:

$$P(\text{like}) = \frac{c(\text{like})}{N}, \text{ where } N \text{ is the number of tokens in the corpus}$$

# Maximum likelihood estimation (MLE)

For a unigram model estimated from a corpus  $w_1 \dots w_k$ , the likelihood function is given by:

$$P(w_1)P(w_2) \dots P(w_k)$$

e.g. for the Bible:

$$P(In)P(the)P(beginning) \dots P(Amen)P(.)$$

If  $N = 790010$  and  $c(In) = 359$  and  $c(the) = 96660$ , etc., then the above expression is **maximized** when

$$P(In) = \frac{359}{790010} \text{ and } P(the) = \frac{96660}{790010} \text{ etc.}$$

Similarly for bigrams, trigrams, etc.

# Counting bigrams

... I would like to know your plans for ...

I would: 1

# Counting bigrams

... I would like to know your plans for ...

I would: 1      would like: 1

# Counting bigrams

... I would like to know your plans for ...

I would: 1      would like: 1      like to: 1

# Bigram probabilities

The word *I* occurred 1000 times...

- ... 20 times, the next word was *like*
- ... 200 times, the next word was *am*
- ... 100 times, the next word was *have*
- etc.

From the counts, we can estimate *bigram probabilities*:

- $P(\textit{like}|\textit{I}) = 0.02$
- $P(\textit{am}|\textit{I}) = 0.2$
- $P(\textit{have}|\textit{I}) = 0.1$
- etc.

# Trigram probabilities

The sequence *I like* occurred 20 times...

- ... 5 times, the next word was *to*
- ... 4 times, the next word was *that*
- ... 1 time, the next word was *apples*
- etc.

From the counts, we can estimate *trigram probabilities*:

- $P(\textit{to}|\textit{I like}) = 0.25$
- $P(\textit{that}|\textit{I like}) = 0.2$
- $P(\textit{apples}|\textit{I like}) = 0.05$
- etc.



# Example

- Corpus with 19 tokens (including punctuation):  
*I live in Boston.*  
*I like ants.*  
*Ants like honey.*  
*Therefore I like honey too.*
- What is  $P(\text{I like Boston})$  using a unigram model based on the above corpus?
- What is  $P(\text{I like honey})$  using a bigram model?
- What is  $P(\text{I like Boston})$  using a bigram model?

# Example

- Corpus with 19 tokens (including punctuation):

*I live in Boston.*

*I like ants.*

*Ants like honey.*

*Therefore I like honey too.*

- What is  $P(\text{I like Boston})$  using a unigram model based on the above corpus?  $(3/19)(3/19)(1/19)$
- What is  $P(\text{I like honey})$  using a bigram model?  
 $(3/19)(2/3)(2/3)$
- What is  $P(\text{I like Boston})$  using a bigram model?  
 $(3/19)(2/3)(0/3) = 0$

# Problems with Maximum likelihood estimation

*Data sparsity* is a problem for the straightforward MLE method.

E.g.  $\frac{c(\text{"I have a unicorn"})}{c(\text{"I have a"})}$

- What if there are no occurrences of “I have a unicorn”?
- What if there are no occurrences of “I have a”?

Regardless of how much data you have, this will happen over and over.

# DD2418 Language Engineering:

## 3c: Zero probabilities, and what to do about them

Johan Boye, KTH

# Zero-probabilities, and what to do about them

A problem with  $n$ -gram models is that many sensible word sequences will have zero-probabilities.

3 techniques to solve this problems:

- Smoothing
- Backoff
- Linear interpolation

# Laplace smoothing

*Smoothing*: Transfer some of the probability mass from the seen sequences to the unseen sequences.

Easiest variant: *Laplace (add-one)* smoothing.

- Previously (Maximum Likelihood Estimation):

$$P_{MLE}(w_i|w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

- Now: (MLE with Laplace smoothing):

$$P_{Laplace}(w_i|w_{i-1}) = \frac{c(w_{i-1}, w_i) + 1}{c(w_{i-1}) + V}$$

where  $V$  is the size of the vocabulary (number of unique words).

# Laplace smoothing (example)

- Corpus with 19 tokens (including punctuation):

*I live in Boston.*

*I like ants.*

*Ants like honey.*

*Therefore I like honey too.*

- What is  $P(\text{I like honey})$  using a bigram model with Laplace smoothing?
- What is  $P(\text{I like Boston})$  using a bigram model with Laplace smoothing?

# Laplace smoothing (example)

- Corpus with 19 tokens (including punctuation):

*I live in Boston.*

*I like ants.*

*Ants like honey.*

*Therefore I like honey too.*

- What is  $P(\text{I like honey})$  using a bigram model with Laplace smoothing?  $(\frac{3+1}{19+10})(\frac{2+1}{3+10})(\frac{2+1}{3+10})$
- What is  $P(\text{I like Boston})$  using a bigram model with Laplace smoothing?  $(\frac{3+1}{19+10})(\frac{2+1}{3+10})(\frac{0+1}{3+10})$



# Advanced smoothing

Laplace smoothing turns out to be too crude for many applications

- ...but it is useful in some machine learning contexts (more on this later in the course)

More sophisticated smoothing methods exist, e.g. the *interpolated Kneser-Ney* method.

We won't cover these in the course.

If a particular  $n$ -gram is not present in the training corpus, then use  $(n - 1)$ -grams instead.

If the  $(n - 1)$ -grams do not exist either, then use  $(n - 2)$ -grams, etc.

Example: Suppose that  $P(\text{ants}|\text{really like}) = 0$  in a trigram model. Then compute the probability as:

$$\hat{P}(\text{ants}|\text{really like}) = P(\text{ants}|\text{like})P(\text{like}|\text{really})$$

Note that this computation will most likely underestimate the actual probability (why?).

# Linear interpolation

Estimate  $\hat{P}(w_i|w_{i-1})$  as

$$\lambda_1 P(w_i|w_{i-1}) + \lambda_2 P(w_i) + \lambda_3$$

where the  $\lambda$ s sum to 1.

$$\sum_i \lambda_i = 1$$

Typically  $\lambda_1 = 0.99$ ,  $\lambda_2 = 0.01 - \lambda_3$ ,  $\lambda_3 = 10^{-6}$

This idea naturally extends to 3-grams, etc.

## Assignment 2, problem 1

Write a program that computes all bigram probabilities from a given (training) corpus, and stores it in a file. Practical issue:

We will use log-probabilities rather than probabilities

- ... using the natural logarithm (because it's simpler)
- $-11.99225$  rather than  $0.0000061919939907$ .

This is to avoid underflow when computing with very small probabilities.

... and we can add rather than multiply

- $\log(p_1 \times \dots \times p_n) = \log(p_1) + \dots + \log(p_n)$

# Probability of a sentence

When calculating the probability of a sentence, it is useful to include punctuation or boundary symbols, e.g.

$$P(\langle b \rangle \text{ I like Boston } \langle b \rangle) = \\ P(\text{I}|\langle b \rangle) \times P(\text{like}|\text{I}) \times P(\text{Boston}|\text{like}) \times P(\langle b \rangle|\text{Boston})$$

- $P(\text{I}|\langle b \rangle)$ : Probability that “I” will be the first word of a sentence.
- $P(\langle b \rangle|\text{Boston})$ : Probability that “Boston” will be the last word of a sentence.

# $n$ -gram models and linguistic structure

How much linguistic structure is captured by  $n$ -gram-models?

- Higher  $n \Rightarrow$  we capture more language structure, BUT
- Higher  $n \Rightarrow$  we need more training data to get accurate probabilities.
- 4-grams and above require Google quantities of data, OR a restricted domain!

Long-distance dependencies will always be a problem (regardless of the choice of  $n$ ):

**The struggle** *between conservatives and socialists is being fought on many fronts.*


# DD2418 Language Engineering: 3d: Evaluation of language models

Johan Boye, KTH

# Evaluation of $n$ -gram models

*Extrinsic evaluation:* Put your  $n$ -gram to use in an application, e.g. a machine translation system. Measure the performance.

*Intrinsic evaluation:* Compute the *entropy* of the model.





# Information

The *information* of an outcome having probability  $p$  is

$$-\log_2 p$$

Information is measured in *bits*.

# Entropy

The *entropy* of a random variable  $X$  is the expected value of the information.

$$H(X) = - \sum_{i=1}^n P(X = x_i) \log_2 P(X = x_i)$$

Special case: We assume that  $0 \log_2 0 = 0$ .

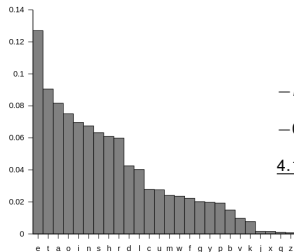
The entropy of  $X$  is a measure of the difficulty of predicting the value of  $X$ .

# Entropy examples

Entropy for a-z using a uniform distribution:

$$-\sum_{i=1}^{26} \frac{1}{26} \log_2 \frac{1}{26} = -\log_2 \frac{1}{26} = \log_2 26 = \underline{4.7}$$

Entropy for a-z using probabilities for English:



$$\begin{aligned} & -P(a) \log_2 P(a) & -P(b) \log_2 P(b) & -\dots & -P(z) \log_2 P(z) = \\ & -0.082 \log_2 0.082 & -0.015 \log_2 0.015 & -\dots & -0.00074 \log_2 0.00074 = \\ & \underline{4.18} \end{aligned}$$

Entropy if  $P(a) = 1$ :  $-1 \log_2 1 - 0 \log_2 0 - \dots 0 \log_2 0 = 0$

# Basic properties of entropy

More predictability  $\Leftrightarrow$  lower entropy.

The entropy of a certain event is 0.

Maximum entropy is obtained if all outcomes are equally probable.

# Entropy and language

What is the entropy at different points in a sentence?

The recent ①

# Entropy and language

What is the entropy at different points in a sentence?

The recent ① results mean that ②

# Entropy and language

What is the entropy at different points in a sentence?

The recent ① results mean that ②  
scientists now need to ③

# Entropy and language

What is the entropy at different points in a sentence?

The recent ① results mean that ②  
scientists now need to ③ go back to  
the drawing ④



# Entropy and language

What is the entropy at different points in a sentence?

The recent ① results mean that ②  
scientists now need to ③ go back to  
the drawing ④ board.

# Cross-entropy

Suppose  $a, b, c, d$  is distributed according to  $p$ :

$$p_a = P(a) = 0.5, p_b = 0.2, p_c = 0.2, p_d = 0.1$$

but we *believe* it is distributed according to  $q$ :

$$q_a = P(a) = 0.1, q_b = 0.2, q_c = 0.2, q_d = 0.5$$

The *cross-entropy of  $p$  on  $q$*  is then computed as:

$$H(X) = - \sum_{i=a,b,c,d} p_i \log_2 q_i$$

Cross-entropy measures how difficult it is to predict the symbol under this belief.

# Entropy as an evaluation metric

Entropy can be used as an evaluation metric for language models.

Given a model  $P$  estimated from a *training corpus*, one can approximate the entropy as:

$$-\frac{1}{N} \log_2 P(w_1, w_2, \dots, w_N)$$

where  $w_1, w_2, \dots, w_N$  is a very long sequence of words from a *test corpus*.

The above computation really approximates the *cross-entropy* of the test set on  $P$ , which is an upper bound of the entropy of  $P$ .

# Entropy as an evaluation metric

The *lower* the entropy of the test corpus, the *better* the language model learned from the training corpus.

The tacit assumption here is that the test corpus is representative of actual data.

## Assignment 2, problem 3

- Write a program that evaluates a language model (a model constructed with your program in (a)) on a given test set.
- Build a number of models and evaluate them on different test sets.
- Draw conclusions.